

# Intro to Data Science - HW 3

Copyright Jeffrey Stanton, Jeffrey Saltz, and Jasmina Tacheva

```
# Enter your name here: Ryan Tervo
# Course Number: IST 687
# Assignment Name: Homework #3
# Due Date: 31 Oct 2022
# Submitted Date: 31 Oct 2022
```

Attribution statement: (choose only one and delete the rest)

```
# 1. I did this homework by myself, with help from the book and the professor.
```

## Reminders of things to practice from last week:

Make a data frame `data.frame( )`  
 Row index of max/min `which.max( )` `which.min( )`  
 Sort value or order rows `sort( )` `order( )`  
 Descriptive statistics `mean( )` `sum( )` `max( )`  
 Conditional statement if (condition) “true stuff” else “false stuff”

## This Week:

Often, when you get a dataset, it is not in the format you want. You can (and should) use code to refine the dataset to become more useful. As Chapter 6 of Introduction to Data Science mentions, this is called “data munging.” In this homework, you will read in a dataset from the web and work on it (in a data frame) to improve its usefulness.

## Part 1: Use `read_csv( )` to read a CSV file from the web into a data frame:

- A. Use R code to read directly from a URL on the web. Store the dataset into a new dataframe, called `dfComps`. The URL is:

“<https://intro-datascience.s3.us-east-2.amazonaws.com/companies1.csv>”

**Hint:** use `read_csv( )`, not `read.csv( )`. This is from the **tidyverse** package. Check the help to compare them.

```
# CODE TO PREVENT WARNINGS/MESSAGES:
#```${r setup, include=FALSE}
#knitr::opts_chunk$set(echo = TRUE, warning = FALSE, message = FALSE)
#```

# IMPORT LIBRARIES:
library(tidyverse)
```

```
## — Attaching packages ————— tidyverse 1.3.2 —
## ggplot2 3.4.0 purrr 0.3.5
```

```
## tibble 3.1.8      dplyr 1.0.10
## tidyr 1.2.1      stringr 1.4.1
## readr 2.1.3      forcats 0.5.2
## — Conflicts ————— tidyverse_conflicts() —
## dplyr::filter() masks stats::filter()
## dplyr::lag() masks stats::lag()
```

```
library(dplyr)
library(stringr)

# DEFINE THE VARIABLES:
fileName <- "https://intro-datascience.s3.us-east-2.amazonaws.com/companies1.csv"

# READ EXCEL FILE USING WEBSITE FILE:
dfComps <- data.frame(read_csv(fileName, show_col_types = F))
dfComps0 <- dfComps
```

## Part 2: Create a new data frame that only contains companies with a homepage URL:

- E. Use **subsetting** to create a new dataframe that contains only the companies with homepage URLs (store that dataframe in **urlComps**).

```
# Use Subsetting to create a new dataframe that contains homepage URLs.
## Each homepage_url will check for either http or www. If those exist then it subsets
it into the urlComps
urlComps <- dfComps[!is.na(dfComps$homepage_url), ]

# Other alternative means to find website entries.
#urlComps <- dfComps[grepl("http", dfComps$homepage_url) | grepl("www", dfComps$homepage_
url), ]
#urlComps <- dfComps[grepl("", dfComps$homepage_url), ]
```

- D. How many companies are missing a homepage URL?

```
# Perform Calculation:
numMissingURL <- nrow(dfComps) - nrow(urlComps)

# Display Results:
printString <- paste('There are ', numMissingURL, ' companies missing there homepage URL.', se
p = "")
print(printString, quote = F)
```

```
## [1] There are 3323 companies missing there homepage URL.
```

## Part 3: Analyze the numeric variables in the dataframe.

- G. How many **numeric variables** does the dataframe have? You can figure that out by looking at the output of **str(urlComps)**.

## H. What is the average number of funding rounds for the companies in `urlComps`?

```
# PART G:
# Perform Calculation:
#str(urlComps) # Initially ran, output examined, and then commented out for readability.

# Display Results:
printString <- paste('Inspecting the structure of urlComps there are 2 numeric variables.', sep = "")
print(printString, quote = F)
```

```
## [1] Inspecting the structure of urlComps there are 2 numeric variables.
```

```
# PART H:
# Perform Calculation:
aveNumFundingRounds <- mean(urlComps$funding_rounds)

# Display Results:
print("", quote = F)
```

```
## [1]
```

```
printString <- paste('The average funding rounds in urlComps is ', round(aveNumFundingRounds, 2), ".", sep = "")
print(printString, quote = F)
```

```
## [1] The average funding rounds in urlComps is 1.73.
```

## I. What year was the oldest company in the dataframe founded?

**Hint:** If you get a value of “NA,” most likely there are missing values in this variable which preclude R from properly calculating the min & max values. You can ignore NAs with basic math calculations. For example, instead of running `mean(urlComps$founded_year)`, something like this will work for determining the average (note that this question needs to use a different function than ‘mean’).

```
#mean(urlComps$founded_year, na.rm=TRUE)
# Example code: mean(urlComps$founded_year, na.rm=TRUE)

# Perform Calculation:
oldestCompYear <- min(urlComps$founded_year, na.rm = TRUE)

# Display Results:
printString = paste('The oldest company was founded in the year: ', oldestCompYear, sep = "")
print(printString, quote = F)
```

```
## [1] The oldest company was founded in the year: 1900
```

```
#####
# ALTERNATIVE CALCULATION:
#####
#   Perform Calculation:
dfComp1 <- dfComps[ , c('name', 'founded_year' )] %>% drop_na()
dfComp1 <- dfComp1[dfComp1$founded_year == min(dfComp1$founded_year) , ]

#   Display Results:
printString = paste('The oldest company was "', dfComp1[1, 1], '" and was founded in the year:
', dfComp1[1, 2], '.', sep = "")
print(printString, quote = F)
```

```
## [1] The oldest company was "The University of Nottingham" and was founded in the year: 1900
.
```

```
#your code goes here
```

## Part 4: Use string operations to clean the data.

- K. The **permalink** variable in **urlComps** contains the name of each company but the names are currently preceded by the prefix “/organization/”. We can use `str_replace()` in tidyverse or `gsub()` to clean the values of this variable:

```
#   Perform Calculation:
urlComps$permalink <- str_replace(urlComps$permalink, "/organization/", "")
```

- L. Can you identify another variable which should be numeric but is currently coded as character? Use the `as.numeric()` function to add a new variable to **urlComps** which contains the values from the char variable as numbers. Do you notice anything about the number of NA values in this new column compared to the original “char” one?

```
#   Perform Calculation:
urlComps$funding_total_usd_num <- as.numeric(urlComps$funding_total_usd)
```

```
## Warning: NAs introduced by coercion
```

```
#   Inspect urlComps
head(urlComps)
```

```
##           permalink           name           homepage_url
## 1           waywire           #waywire http://www.waywire.com
## 2 tv-communications &TV Communications http://enjoyandtv.com
## 3   rock-your-paper 'Rock' Your Paper http://www.rockyourpaper.org
## 4 in-touch-network (In)Touch Network http://www.InTouchNetwork.com
## 5             n-plusn             #NAME? http://plusn.com
## 7     club-domains     .Club Domains http://nic.club/
##
##                                     category_list
## 1                                     |Entertainment|Politics|Social Media|News|
## 2                                     |Games|
```

```
## 3 |Publishing|Education|
## 4 |Electronics|Guides|Coffee|Restaurants|Music|iPhone|Apps|Mobile|iOS|E-Commerce|
## 5 |Software|
## 7 |Software|

##      market funding_total_usd      status country_code state_code
## 1      News      1 750 000  acquired      USA      NY
## 2      Games      4 000 000  operating      USA      CA
## 3  Publishing      40 000  operating      EST      <NA>
## 4 Electronics      1 500 000  operating      GBR      <NA>
## 5      Software      1 200 000  operating      USA      NY
## 7      Software      7 000 000      <NA>      USA      FL

##      region      city funding_rounds founded_at founded_month
## 1 New York City New York      1      1/6/12      2012-06
## 2 Los Angeles Los Angeles      2      <NA>      <NA>
## 3 Tallinn Tallinn      1 26/10/2012      2012-10
## 4 London London      1      1/4/11      2011-04
## 5 New York City New York      2      1/1/12      2012-01
## 7 Ft. Lauderdale Oakland Park      1      10/10/11      2011-10

## founded_quarter founded_year first_funding_at last_funding_at
## 1      2012-Q2      2012      30/06/2012      30/06/2012
## 2      <NA>      NA      4/6/10      23/09/2010
## 3      2012-Q4      2012      9/8/12      9/8/12
## 4      2011-Q2      2011      1/4/11      1/4/11
## 5      2012-Q1      2012      29/08/2012      4/9/14
## 7      2011-Q4      2011      31/05/2013      31/05/2013

## funding_total_usd_num
## 1      NA
## 2      NA
## 3      NA
## 4      NA
## 5      NA
## 7      NA
```

```
#urlComps[ , c(funding_total_usd)]#, funding_total_usd_num)])

#urlComps$funding_total_usd_num <- as.numeric(str_replace_all(urlComps$funding_total_usd, "[[:
space:]]", ""))

# Inspection:
printString1 <- "It appears that all of the new FUNDING_TOTAL_USD values are NA."
printString2 <- "Upon further inspection it looks like there are spaces in the FUNDING_TOTAL_U
SD values which are preventing the numeric conversion."

# Display Output:
print(printString1, quote = F)
```

```
## [1] It appears that all of the new FUNDING_TOTAL_USD values are NA.
```

```
print(printString2, quote = F)
```

```
## [1] Upon further inspection it looks like there are spaces in the FUNDING_TOTAL_USD values
which are preventing the numeric conversion.
```

M. To ensure the char values are converted correctly, we first need to remove the spaces between the digits in the variable. Check if this works, and explain what it is doing:

```
library(stringi)
urlComps$funding_new <- stri_replace_all_charclass(urlComps$funding_total_usd, "\\p{WHITE_SPACE}+", "")

# Perform Calculation:
# In order to convert the dollar values to numbers the spaces need to be removed.
# METHOD #1
library(stringi)
urlComps$funding_new <- stri_replace_all_charclass(urlComps$funding_total_usd, "\\p{WHITE_SPACE}+", "")

# METHOD #2:
dfTemp <- str_replace_all(urlComps$funding_total_usd, "[[:space:]]+", "")

# DISPLAY RESULTS:
head(dfTemp)
```

```
## [1] "1750000" "4000000" "40000" "1500000" "1200000" "7000000"
```

```
head(urlComps$funding_new)
```

```
## [1] "1750000" "4000000" "40000" "1500000" "1200000" "7000000"
```

```
printString <- 'The "stri_replace_all_charclass" removes all types of the white space in the s
tring regardless of position.'
print(printString, quote = F)
```

```
## [1] The "stri_replace_all_charclass" removes all types of the white space in the string reg
ardless of position.
```

```
Error in stri_replace_all_charclass(urlComps$funding_total_usd, "\\p{WHITE_SPACE}+", : object '
urlComps' not found
Traceback:

1. stri_replace_all_charclass(urlComps$funding_total_usd, "\\p{WHITE_SPACE}+",
.  )
```

N. You are now ready to convert **urlComps\$funding\_new** to numeric using **as.numeric()**.

Calculate the average funding amount for **urlComps**. If you get “NA,” try using the **na.rm=TRUE** argument from problem I.

```
#   Perform the Calculation:
urlComps$funding_new <- as.numeric(urlComps$funding_new)
```

```
## Warning: NAs introduced by coercion
```

```
meanFunding <- mean(urlComps$funding_new, na.rm = TRUE)

printMeanFundingAmount <- format(round(meanFunding, 2), nsmall = 2, big.mark = ",")
#   Take the mean funding value and format for better display.

#   Display Output:
printString <- paste('The mean funding amount was $', printMeanFundingAmount, ".", sep = "")
print(printString, quote = F)
```

```
## [1] The mean funding amount was $18,321,551.47.
```

Sample three unique observations from `urlComps$funding_rounds`, store the results in the vector 'observations'

```
#   Perform Calculation:
numObservations = 3
observations <- sample(urlComps$funding_rounds, numObservations, replace = FALSE)

#   Display Output:
print(observations, quote = FALSE)
```

```
## [1] 1 3 1
```

Take the mean of those observations

```
#   Perform Calculation:
meanObservations <- round(mean(observations), 2)

#   Display Output:
printString <- paste('The mean observations is: ', meanObservations, sep = '')
print(printString, quote = FALSE)
```

```
## [1] The mean observations is: 1.67
```

Do the two steps (sampling and taking the mean) in one line of code

```
#   Perform Calculation:
meanObservations = round(mean(sample(urlComps$funding_rounds, numObservations, replace = FALSE
)), 2)

#   Display Output:
printString <- paste('The mean observations is: ', meanObservations, sep = '')
print(printString, quote = FALSE)
```

```
## [1] The mean observations is: 2.67
```

Explain why the two means are (or might be) different

Use the `replicate()` function to repeat your sampling of three observations of `urlComps$funding_rounds` observations five times. The first argument to `replicate()` is the number of repeats you want. The second argument is the little chunk of code you want repeated.

```
# Define Variables:
numSamples <- 3
numTrials  <- 5

# Perform Calculations:
sampleMeans <- replicate(numTrials, mean(sample(urlComps$funding_rounds, numSamples, replace =
  FALSE)), simplify = TRUE)
sampleMeans <- round(sampleMeans, 2)

# Display Results:
print(sampleMeans)
```

```
## [1] 1.67 1.00 1.33 1.00 2.67
```

Rerun your replication, this time doing 20 replications and storing the output of `replicate()` in a variable called **values**.

```
# Define Variables:
numSamples <- 3
numTrials  <- 20

# Perform Calculations:
values <- replicate(numTrials, mean(sample(urlComps$funding_rounds, numSamples, replace = FALSE),
  simplify = TRUE)
values <- round(values, 2)

# Display Results:
print(values)
```


```
## [1] 2.00 1.00 2.00 1.67 1.00 1.00 1.67 2.00 1.67 2.33 1.67 2.33 3.67 1.33 2.33
## [16] 1.33 2.00 1.67 1.00 3.00
```

Generate a **histogram** of the means stored in **values**.

```
hist(values)
```



## Histogram of values




Rerun your replication, this time doing 1000 replications and storing the output of `replicate()` in a variable called **values**, and then generate a histogram of **values**.

```
# Define Variables:
numSamples <- 3
numTrials  <- 1000

# Perform Calculations:
values <- replicate(numTrials, mean(sample(urlComps$funding_rounds, numSamples, replace = FALSE), simplify = TRUE))

# Display Results:
hist(values)
```

## Histogram of values




Repeat the replicated sampling, but this time, raise your sample size from 3 to 22. How does that affect your histogram? Explain in a comment.

```
# Define Variables:
numSamples <- 22
numTrials <- 1000

# Perform Calculations:
values <- replicate(numTrials, mean(sample(urlComps$funding_rounds, numSamples, replace = FALSE), simplify = TRUE))
#values <- round(values, 2)

# Display Results:
#print(values)
hist(values)
```

## Histogram of values



Explain in a comment below, the last three histograms, why do they look different?

```
printString1 = 'As the number of samples and trials increase then the sample mean distribution
approaches a normal distribution.'
printString2 = 'Initially there were too few samples and trials conducted so the histogram did
not show a normal distribution.'

print(printString1, quote = FALSE)
```

```
## [1] As the number of samples and trials increase then the sample mean distribution approach
es a normal distribution.
```

```
print(printString2, quote = FALSE)
```

```
## [1] Initially there were too few samples and trials conducted so the histogram did not show
a normal distribution.
```