

# Intro to Data Science - HW 5

Copyright Jeffrey Stanton, Jeffrey Saltz, and Jasmina Tacheva

```
# Enter your name here: Ryan Tervo
# Course Number: IST 687
# Assignment Name: Homework #5
# Due Date: 14 Nov 2022
# Submitted Date: 14 Nov 2022
```

Attribution statement: (choose only one and delete the rest)

```
# 1. I did this homework by myself, with help from the book and the professor.
```

**This module: Data visualization** is important because many people can make sense of data more easily when it is presented in graphic form. As a data scientist, you will have to present complex data to decision makers in a form that makes the data interpretable for them. From your experience with Excel and other tools, you know that there are a variety of **common data visualizations** (e.g., pie charts). How many of them can you name?

The most powerful tool for data visualization in R is called **ggplot**. Written by computer/data scientist **Hadley Wickham**, this “**graphics grammar**” tool builds visualizations in layers. This method provides immense flexibility, but takes a bit of practice to master.

## Step 1: Make a copy of the data

- Read the **who** dataset from this URL: <https://intro-datascience.s3.us-east-2.amazonaws.com/who.csv> into a new dataframe called **tb**.

Your new dataframe, **tb**, contains a so-called **multivariate time series**: a sequence of measurements on 23 Tuberculosis-related (TB) variables captured repeatedly over time (1980-2013). Familiarize yourself with the nature of the 23 variables by consulting the dataset’s codebook which can be found here: [https://intro-datascience.s3.us-east-2.amazonaws.com/TB\\_data\\_dictionary\\_2021-02-06.csv](https://intro-datascience.s3.us-east-2.amazonaws.com/TB_data_dictionary_2021-02-06.csv).

```
# IMPORT LIBRARIES:
library(tidyverse)
```

```
## — Attaching packages ————— tidyverse 1.3.2 —
##  ggplot2 3.4.0    purrr  0.3.5
##  tibble  3.1.8    dplyr  1.0.10
##  tidyr   1.2.1    stringr 1.4.1
##  readr   2.1.3    forcats 0.5.2
## — Conflicts ————— tidyverse_conflicts() —
##  dplyr::filter() masks stats::filter()
##  dplyr::lag()    masks stats::lag()
```

```
library(dplyr)
library(stringr)
```

```
#   DEFINE THE VARIABLES:
fileName <- "https://intro-datascience.s3.us-east-2.amazonaws.com/who.csv"

#   READ EXCEL FILE USING WEBSITE FILE:
tb <- data.frame(read_csv(fileName)) # show_col_types = FALSE))
```

```
## Rows: 5769 Columns: 23
## — Column specification —————
## Delimiter: ","
## chr (1): iso2
## dbl (22): year, new_sp, new_sp_m04, new_sp_m514, new_sp_m014, new_sp_m1524, ...
##
## i Use `spec()` to retrieve the full column specification for this data.
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
tb0 <- tb
```

- B. How often were these measurements taken (in other words, at what frequency were the variables measured)? Put your answer in a comment.

```
#   The measurements were taken annually.
#   This can be seen in the tb$year column in which data is collected on a yearly or annual basis
.
```

## Step 2: Clean-up the NAs and create a subset

- A. Let's clean up the iso2 attribute in **tb**

Hint: use *is.na()* – well use *! is.na()*

```
tb1 <- tb[!is.na(tb$iso2), ]
```

- B. Create a subset of **tb** containing **only the records for Canada (“CA” in the iso2 variable)**. Save it in a new dataframe called **tbCan**. Make sure this new df has **29 observations and 23 variables**.

```
#   Create tb2 per the instructions: only CA
tbCan <- tb1[tb1$iso2 == "CA", ]

#   Verify tb2 has Correct Dimensions:
numRow <- nrow(tbCan)
numCol <- ncol(tbCan)

#   Display the Results:
printString = paste('Dataframe tb2 has ', numRow, ' rows (observations) and ', numCol, ' columns (variables).', sep = '')
print(printString, quote = FALSE)
```

```
## [1] Dataframe tb2 has 29 rows (observations) and 23 columns (variables).
```

- C. A simple method for dealing with small amounts of **missing data** in a numeric variable is to **substitute the mean of**

the variable in place of each missing datum.

This expression locates (and reports to the console) all the missing data elements in the variable measuring the **number of positive pulmonary smear tests for male children 0-4 years old** (there are 26 data points missing)

```
tbCan$new_sp_m04[is.na(tbCan$new_sp_m04)]
```

```
## [1] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
## [26] NA
```

```
Error in eval(expr, envir, enclos): object 'tbCan' not found
Traceback:
```

D. Write a comment describing how that statement works.

```
# Understanding how the statement works is helpful to look at the three individual parts
# Each part has an output.
# Part 2 and part 3 have an input.

#           tbCan$new_sp_m04[is.na(tbCan$new_sp_m04)]
#part           3           2           1

# Part 1:  tbCan$new_sp_m04
#           input:
#           ouput:  a vector of the tbCan$new_sp_mo4 column values.

# Part 2:  is.na('Part 1')
#           input:  A column of values from tbCan$new_sp_m04
#           output: A column of values TRUE and FALSE.
#           # For each input element that is 'na' the corresponding output element is TRUE
#           # For each input element that is not 'na' then the corresponding output element
is FALSE.
#           # In this case the only True elements correspond to elements which have NA, per
the test.

# Part 3:  tbCan$new_sp_m04['Part2']
#           input:  A vector of TRUE and FALSE values which has a length as the number of el
ements in tbCan$new_sp_m04
#           output: A vector of values in "tbCan$new_sp_m04" of only the associated TRUE ele
ments in Part 2.

# FINAL OUTPUT: A vector, elements are all na, the length is equal to the number of na's in
tbCan$new_sp_m04 column.
```

E. Write 4 more statements to check if there is missing data for the number of positive pulmonary smear tests for: **male and female** children 0-14 years old (**new\_sp\_m014** and **new\_sp\_f014**), and **male and female citizens 65 years of age and older**, respectively. What does empty output suggest about the number of missing observations?

```
# Perform Tests
# Creates a vector of TRUE and FALSE for each column based on whether or not it is na.
# SUMs the TRUE and FALSE statements to get total TRUE for each one. TRUE = 1 and FALS
E = 0.
```

```
test0 <- is.na(tbCan$new_sp_m04)
test0Results <- sum(test0)

test1 <- is.na(tbCan$new_sp_m014)
test1Results <- sum(test1)

test2 <- is.na(tbCan$new_sp_f014)
test2Results <- sum(test2)

test3 <- is.na(tbCan$new_sp_m65)
test3Results <- sum(test3)

test4 <- is.na(tbCan$new_sp_f65)
test4Results <- sum(test4)

#   Display Output:

printString0 <- paste('The column male      0 - 04      had ', test0Results, ' missing data.',
  sep = '')
printString1 <- paste('The column male      0 - 14      had ', test1Results, ' missing data.',
  sep = '')
printString2 <- paste('The column female    0 - 14      had ', test2Results, ' missing data.',
  sep = '')
printString3 <- paste('The column male      65 and older had ', test3Results, ' missing data.',
  sep = '')
printString4 <- paste('The column female    65 and older had ', test4Results, ' missing data.',
  sep = '')

print(printString0, quote = FALSE)
```

```
## [1] The column male      0 - 04      had 26 missing data.
```

```
print(printString1, quote = FALSE)
```

```
## [1] The column male      0 - 14      had 0 missing data.
```

```
print(printString2, quote = FALSE)
```

```
## [1] The column female    0 - 14      had 0 missing data.
```

```
print(printString3, quote = FALSE)
```

```
## [1] The column male      65 and older had 0 missing data.
```

```
print(printString4, quote = FALSE)
```

```
## [1] The column female    65 and older had 0 missing data.
```

```
# What does empty output suggest about the number of missing observations?
# This suggests that there are no missing fields or no na's.
```

There is an R package called **imputeTS** specifically designed to repair missing values in time series data. We will use this instead of mean substitution.

The **na\_interpolation()** function in this package takes advantage of a unique characteristic of time series data: **neighboring points in time can be used to “guess” about a missing value in between.**

F. Install the **imputeTS** package (if needed) and use **na\_interpolation()** on the variable from part C. Don't forget that you need to save the results back to the **tbCan** dataframe. Also update any attribute discussed in part E (if needed).

```
library(imputeTS)
```

```
## Registered S3 method overwritten by 'quantmod':
##      method      from
## as.zoo.data.frame zoo
```

```
# COMPLETE PART F USING 'na_interpolation()' function.
tbCan$new_sp_m04 <- na_interpolation(tbCan$new_sp_m04)
```

G. Rerun the code from C and E above to check that all missing data have been fixed.

```
# Verify tbCan "na's" have been resolved.
test5 <- is.na(tbCan$new_sp_m04)
test5Results <- sum(test5)
#printString5 <- paste('The column had ', test5Results, ' missing data.', sep = '')
#print(printString5, quote = FALSE)

# RERUN PART C:
tbCan$new_sp_m04[is.na(tbCan$new_sp_m04)]
```

```
## numeric(0)
```

```
# RERUN PART E:
test0 <- is.na(tbCan$new_sp_m04)
test0Results <- sum(test0)

test1 <- is.na(tbCan$new_sp_m014)
test1Results <- sum(test1)

test2 <- is.na(tbCan$new_sp_f014)
test2Results <- sum(test2)

test3 <- is.na(tbCan$new_sp_m65)
test3Results <- sum(test3)

test4 <- is.na(tbCan$new_sp_f65)
test4Results <- sum(test4)
```

```
# Display Output:
```

```
printString0 <- paste('The column male    0 - 04      had ', test0Results, ' missing data.',
  sep = '')
printString1 <- paste('The column male    0 - 14      had ', test1Results, ' missing data.',
  sep = '')
printString2 <- paste('The column female  0 - 14      had ', test2Results, ' missing data.',
  sep = '')
printString3 <- paste('The column male    65 and older had ', test3Results, ' missing data.',
  sep = '')
printString4 <- paste('The column female  65 and older had ', test4Results, ' missing data.',
  sep = '')

print(printString0, quote = FALSE)
```

```
## [1] The column male    0 - 04      had 0 missing data.
```

```
print(printString1, quote = FALSE)
```

```
## [1] The column male    0 - 14      had 0 missing data.
```

```
print(printString2, quote = FALSE)
```

```
## [1] The column female  0 - 14      had 0 missing data.
```

```
print(printString3, quote = FALSE)
```

```
## [1] The column male    65 and older had 0 missing data.
```

```
print(printString4, quote = FALSE)
```

```
## [1] The column female  65 and older had 0 missing data.
```

## Step 3: Use ggplot to explore the distribution of each variable

**Don't forget to install and library the ggplot2 package.** Then:

H. Create a histogram for **new\_sp\_m014**. Be sure to add a title and briefly describe what the histogram means in a comment.

```
# INSTALL LIBRARY:
```


```
library(ggplot2)
```

```
# CREATE PLOT USING GGLOT:
```

```
plot1 <- ggplot(data = tbCan, aes(x = new_sp_m014)) + geom_histogram()
```

```
#   DISPLAY PLOT:
plot1
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



- I. Create histograms (using ggplot) of each of the other three variables from E with ggplot( ). Which parameter do you need to adjust to make the other histograms look right?

```
#   CREATE HISTOGRAMS USING GGLOT:
plot2 <- ggplot(data = tbCan, aes(x = new_sp_f014)) + geom_histogram()
plot2
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
plot3 <- ggplot(data = tbCan, aes(x = new_sp_m65)) + geom_histogram()
plot3
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```





```
plot4 <- ggplot(data = tbCan, aes(x = new_sp_f65)) + geom_histogram()
plot4
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
# Which parameter do you need to adjust to make the other histograms look right?
# The parameter that needs to be adjusted is the number of bins.
```

## Step 4: Explore how the data changes over time

J. These data were collected in a period of several decades (1980-2013). You can thus observe changes over time with the help of a line chart. Create a **line chart**, with **year** on the X-axis and **new\_sp\_m014** on the Y-axis.

```
# CREATE GG PLOT:
plot_J <- ggplot(tbCan, aes(x = year)) + geom_line(aes(y = new_sp_m014))

# DISPLAY PLOT:
plot_J
```




K. Next, create similar graphs for each of the other three variables. Change the **color** of the line plots (any color you want).


```
# CREATE GGLOT: COMBINED
plot_K <- ggplot(tbCan, aes(x = year)) + geom_line(aes(y = new_sp_m014, colour = "blue")) +
  geom_line(aes(y = new_sp_f014, colour = "orange")) +
  geom_line(aes(y = new_sp_m65, colour = "red")) +
  geom_line(aes(y = new_sp_f65, colour = "green"))

# CREATE GGLOT SEPARATE:
plot_K1 <- ggplot(tbCan, aes(x = year)) + geom_line(aes(y = new_sp_m014, colour = "blue"))
plot_K2 <- ggplot(tbCan, aes(x = year)) + geom_line(aes(y = new_sp_f014, colour = "blue"))
plot_K3 <- ggplot(tbCan, aes(x = year)) + geom_line(aes(y = new_sp_m65, colour = "green"))
plot_K4 <- ggplot(tbCan, aes(x = year)) + geom_line(aes(y = new_sp_f65, colour = "green"))


# DISPLAY PLOTS:
plot_K
```

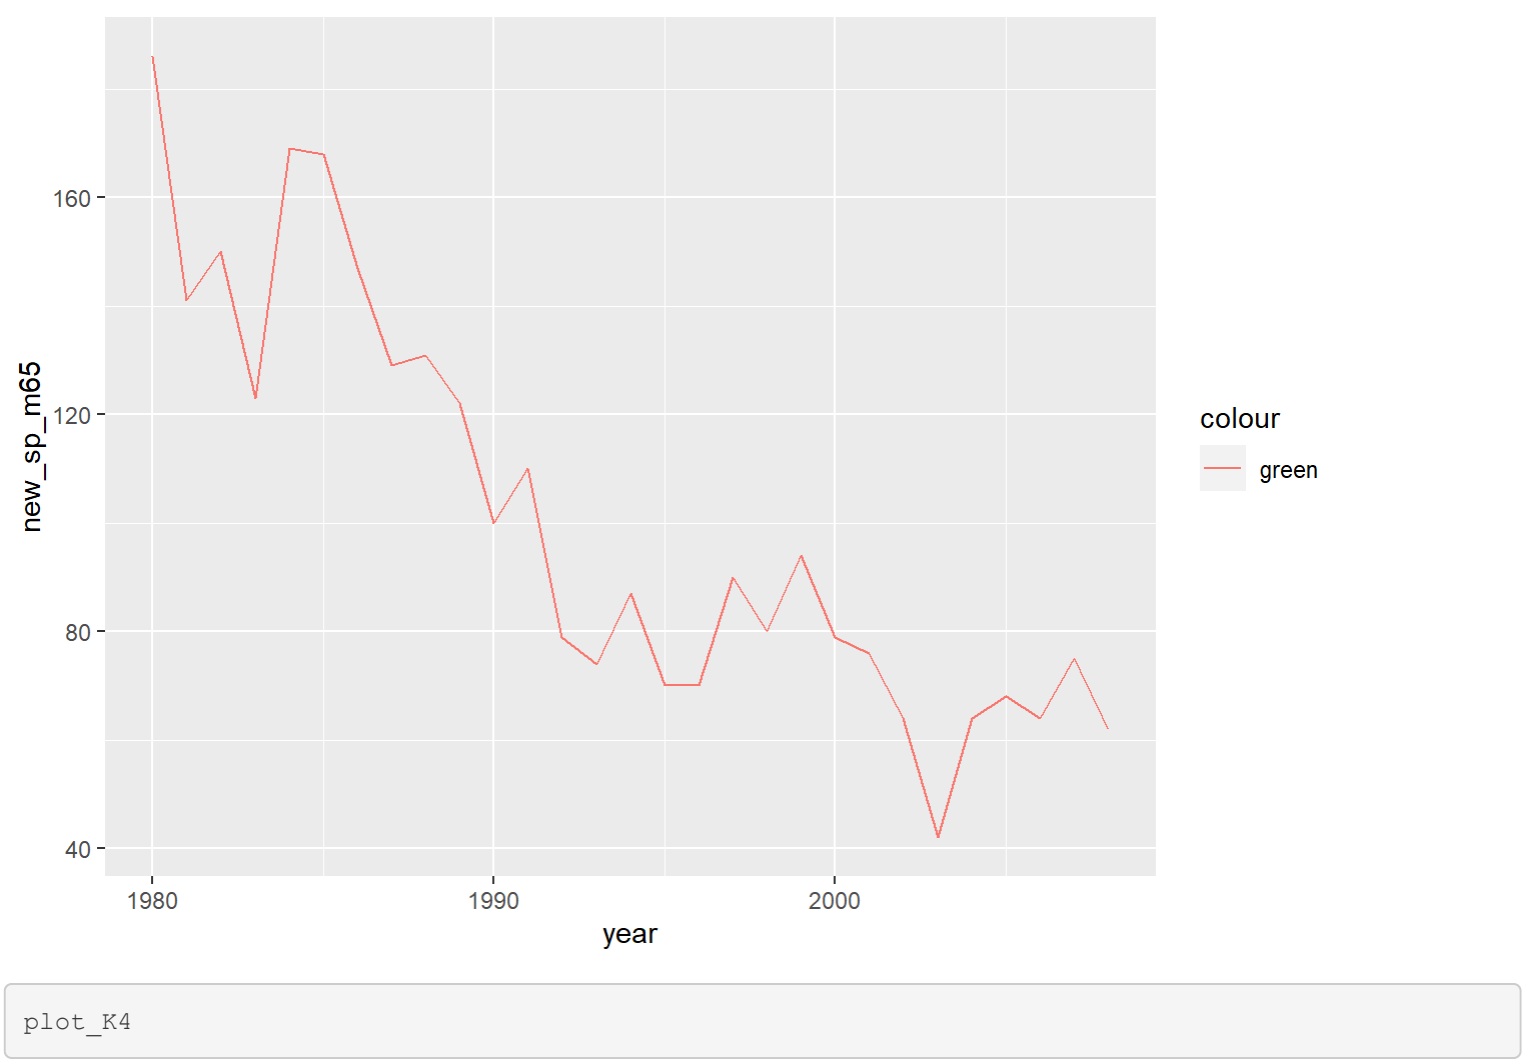


plot\_K1



plot\_K2







- L. Using vector math, create a new variable by combining the numbers from **new\_sp\_m014** and **new\_sp\_f014**. Save the resulting vector as a new variable in the **tbCan** df called **new\_sp\_combined014**. This new variable represents the number of positive pulmonary smear tests for male AND female children between the ages of 0 and 14 years of age. Do the same for SP tests among citizens 65 years of age and older and save the resulting vector in the **tbCan** variable called **new\_sp\_combined65**.

```
# CREATE NEW COLUMNS IN tbCan
tbCan$new_sp_combined014 <- tbCan$new_sp_m014 + tbCan$new_sp_f014
tbCan$new_sp_combined65 <- tbCan$new_sp_m65 + tbCan$new_sp_f65

head(tbCan)
```

```
##      iso2 year new_sp new_sp_m04 new_sp_m514 new_sp_m014 new_sp_m1524
## 872   CA 1980   951         1         NA         12         54
## 873   CA 1981   803         1         NA          8         49
## 874   CA 1982   812         1         NA          6         52
## 875   CA 1983   771         1         NA          9         47
## 876   CA 1984   811         1         NA          3         44
## 877   CA 1985   791         1         NA         11         42
##      new_sp_m2534 new_sp_m3544 new_sp_m4554 new_sp_m5564 new_sp_m65 new_sp_mu
## 872          75          83          100          108          186         NA
## 873          61          64           87          103          141         NA
## 874          66          69           90           91          150         NA
## 875          63          62           90           92          123         NA
## 876          75          58           68           83          169         NA
```




```
## 877          70          59          77          81          168          NA
##      new_sp_f04 new_sp_f514 new_sp_f014 new_sp_f1524 new_sp_f2534 new_sp_f3544
## 872          NA          NA          18          62          51          34
## 873          NA          NA          6          46          57          26
## 874          NA          NA          7          51          57          30
## 875          NA          NA          11         50          50          29
## 876          NA          NA          9          51          59          28
## 877          NA          NA          5          30          56          19
##      new_sp_f4554 new_sp_f5564 new_sp_f65 new_sp_fu new_sp_combined014
## 872          31          33          104          NA          30
## 873          28          35          92          NA          14
## 874          25          38          80          NA          13
## 875          24          35          86          NA          20
## 876          28          36          100         NA          12
## 877          28          48          97          NA          16
##      new_sp_combined65
## 872          290
## 873          233
## 874          230
## 875          209
## 876          269
## 877          265
```

M. Finally, create a **scatter plot**, showing **new\_sp\_combined014** on the x axis, **new\_sp\_combined65** on the y axis, and having the **color and size** of the point represent **year**.

```
#      CREATE GGLOT SEPARATE:
plot_M <- ggplot(tbCan, aes(x = new_sp_combined014, y = new_sp_combined65)) + geom_point(aes(
  size = year, colour = year))

#      DISPLAY PLOTS:
plot_M
```



N. Interpret this visualization – what insight does it provide?

```
# INTERPRETATION:
# At all times there were more cases in the older population than in the younger population.
# There were more cases in the 2000's than in the 1980's.
```